




# Socket Programming and Data Encoding

Peter Sjödin



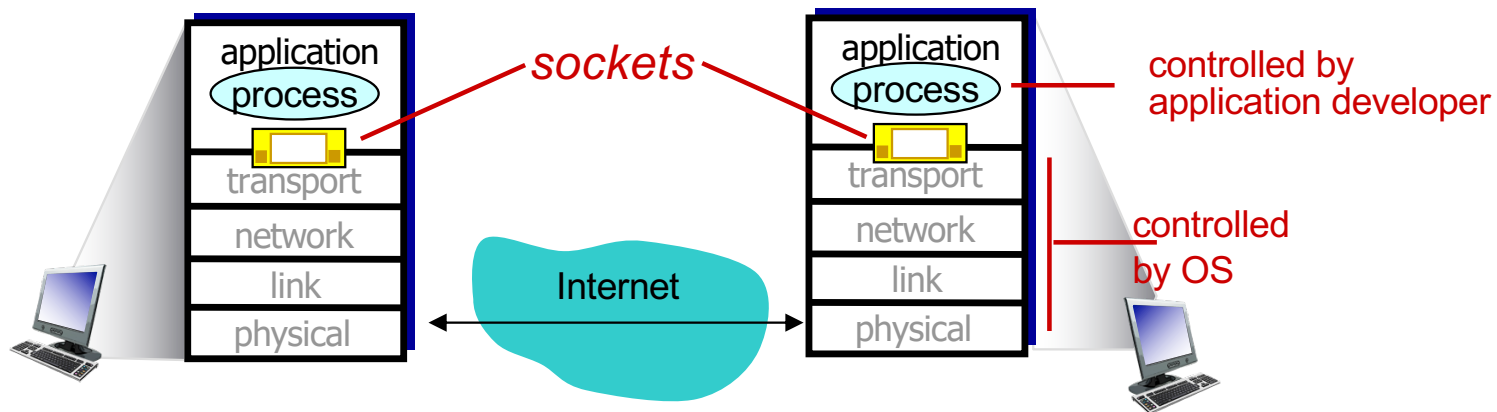


# Synopsis

- The Sockets Applications Programming Interface (API)
- Socket programming in Java
- Java I/O
- Data encoding and decoding

# Sockets

- process sends/receives messages to/from its **socket**
  - How application protocols access transport protocol services





# Socket Programming

- Two socket types for two transport services:
- **TCP**: reliable, byte stream-oriented
  - > *TCP responsible for dividing the stream into packets*
- **UDP**: unreliable, byte block-oriented (datagram)
  - > *Application responsible for dividing data into packets*

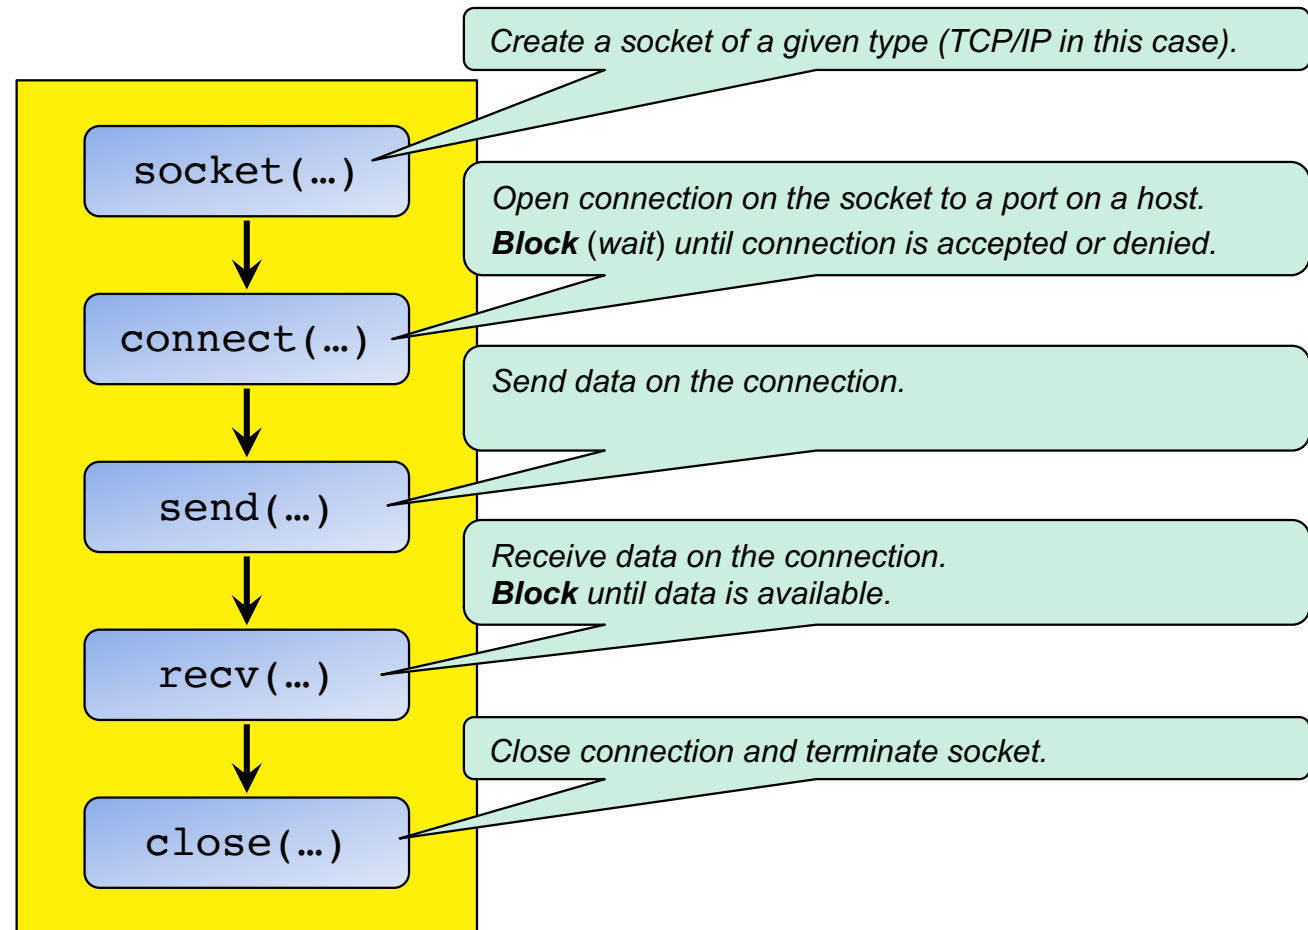


# Socket Programming with TCP

- First, client must open a **TCP connection** to the server
- TCP connection
  - Ordered delivery of data as a *stream of bytes*
    - > Bytes delivered in the same order they are sent
  - Reliable
    - > If the data does not make it to the other side, the application will learn
  - Bi-directional
    - > Both sides (client and server) can send and receive
- When communication is completed, the connection is closed
  - Both client and server may take initiative to close connection (FIN flag)
  - The two directions are closed independently
    - > *A connection can be half open*

# TCP Client

Client creates a socket and performs a number of **socket system calls** on it – functions in the operating system





# Java Streams

- In Java, streams (InputStream and OutputStream) are the basic classes for byte I/O

Write/send all bytes  
in buffer.

**OutputStream()**

`void write(byte [] buffer)`

`void write(byte [] buffer, int offset, int length)`

`void write(int)`

Write/send length bytes from buffer,  
starting at offset.

Write/send a single byte (0 – 255). Larger  
integers will be truncated.



# Java Streams

- Reads data into an existing (pre-allocated) array
- Returns the number of bytes being read
- Returns -1 at end of data
  - End of file, connection closed, etc.

Read/receive bytes into buffer

**InputStream( )**

`int read(byte [] buffer)`

`int read(byte [] buffer, int offset, int length)`

Read/receive at most length bytes into buffer, starting at offset.





# Socket Programming Example Application

1. Client reads data from system input and sends data to server
2. Server receives the data and computes a response
3. Server sends response to client
4. Client receives response and writes it to system output

# Example TCP Client in Java

```
import java.net.*;
class TCPClient {
    private static int BUFFERSIZE=1024;

    public static void main(String argv[]) throws Exception
    {
        // Pre-allocate byte buffers for reading/receiving
        byte[] fromUserBuffer = new byte[BUFFERSIZE];
        byte[] fromServerBuffer = new byte[BUFFERSIZE];

        Socket clientSocket = new Socket("hostname", 6789);

        int fromUserLength = System.in.read(fromUserBuffer); // User input
        clientSocket.getOutputStream().write(fromUserBuffer, 0, fromUserLength);

        int fromServerLength = clientSocket.getInputStream().read(fromServerBuffer);
        System.out.print("FROM SERVER: "); // Use print method since it is a string
        System.out.write(fromServerBuffer, 0, fromServerLength);
        clientSocket.close();
    }
}
```

Defines Socket class

Calls socket() and then connect() system calls to open connection to server "hostname" at port 6789

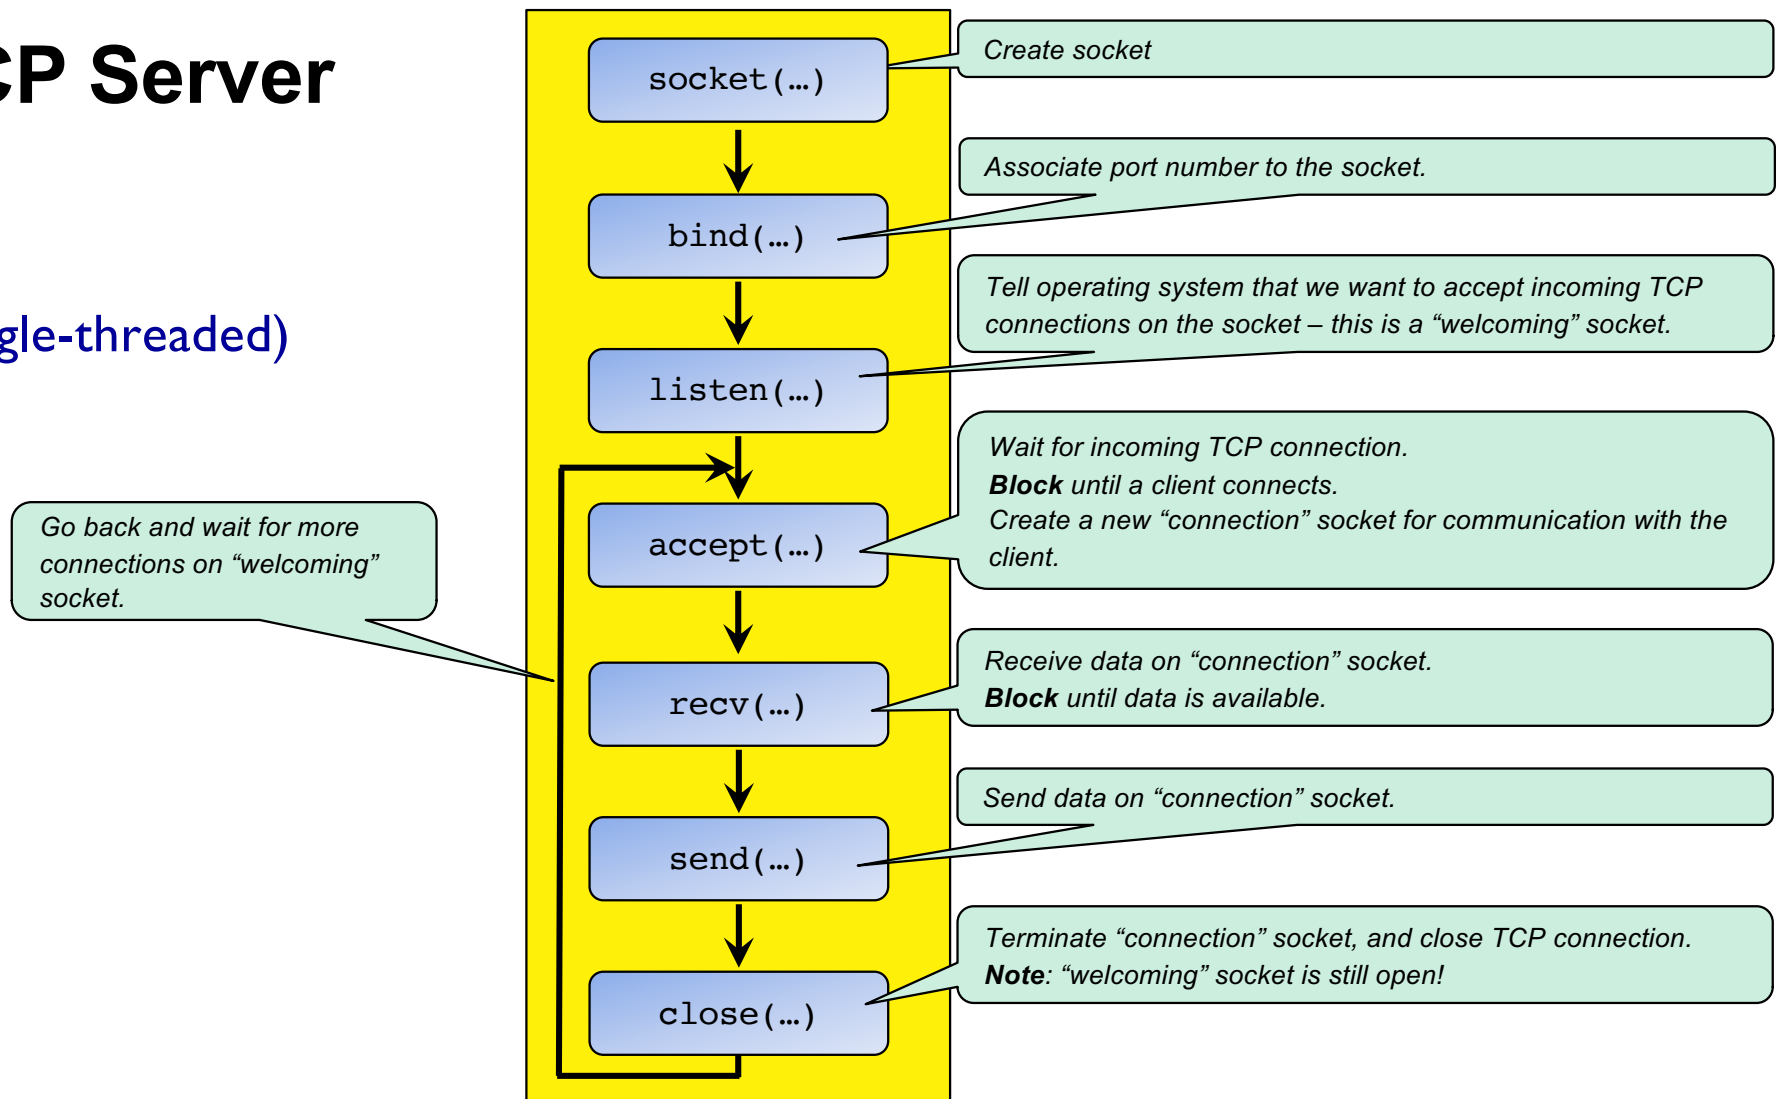
Send bytes on socket

Receive bytes on socket



# TCP Server

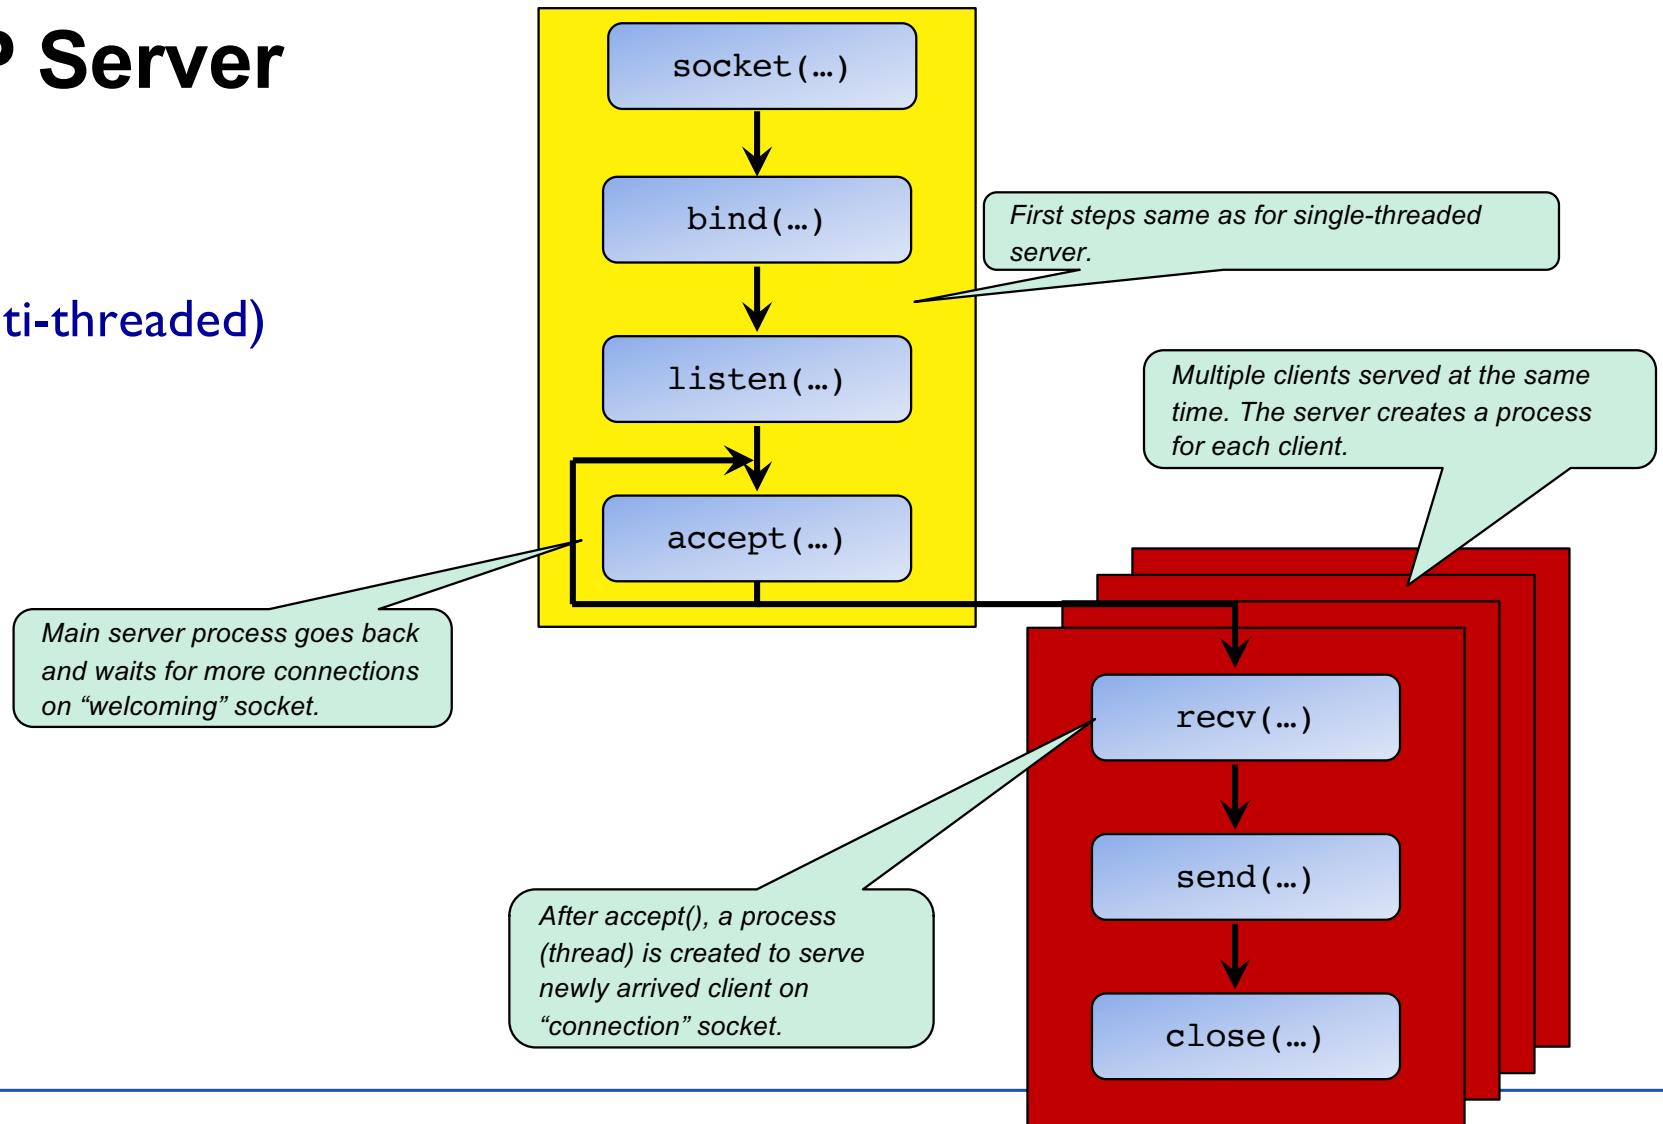
## Sequential (single-threaded) version





# TCP Server

## Concurrent (multi-threaded) version





# TCP Server

```
import java.net.*;
class TCPServer {
    static int BUFFERSIZE=1024;

    public static void main(String argv[]) throws Exception
    {
        ServerSocket welcomeSocket = new ServerSocket(6789);

        while(true) {
            Socket connectionSocket = welcomeSocket.accept();
            byte[] fromClientBuffer = new byte[BUFFERSIZE];
            int fromClientLength = connectionSocket.getInputStream().read(fromClientBuffer);

            // Compute response for client
            byte[] toClientBuffer = ...

            connectionSocket.getOutputStream().write(toClientBuffer);
            connectionSocket.close();
        }
    }
}
```

Calls socket(), bind() and listen() system calls to create a “welcoming” socket with port number 6789.

Wait for a client to connect. Create a new socket for communication with the client.

End of loop statement. Go back and wait for another client.

Sequential or concurrent server?



# Encoding/Decoding

- Programming languages have data types
  - > *Integers, floating point numbers, strings, booleans, etc.*
- To transfer a string between two programs (processes) over the network, for example, we must decide how to represent the string as a sequence of bytes
  - In other words, how convert between bytes and data types?
  - That is what encoding/decoding is about
- Another example: Save a floating point number to file



# Text Encoding/Decoding

- Translate between string symbols (such as "A", "?", "d", "ü", "戏") and bytes
- Define how each symbol is represented as one or more bytes
  - Character encoding standard
  - ASCII, UTF-8, ISO 8859-1, ...



# ASCII

- American Standard Code for Information Interchange
- 7 bits per symbol
  - > *In practice, one byte per symbol*
  - > *8th bit always zero*
  - "A" is ASCII 65 (41 hexadecimal)
  - "z" is ASCII 122 (7A hexadecimal)
  - ASCII 0 to 31 represent control characters
    - > *7 (07 hex) is bell ("beep")*
    - > *10 (0A hex) is line feed*
    - > *13 (0D hex) is carriage return*

	0	1	2	3	4	5	6	7
0	NUL	DLE	space	0	@	P	`	p
1	SOH	DC1 XON	!	1	A	Q	a	q
2	STX	DC2	"	2	B	R	b	r
3	ETX	DC3 XOFF	#	3	C	S	c	s
4	EOT	DC4	\$	4	D	T	d	t
5	ENQ	NAK	%	5	E	U	e	u
6	ACK	SYN	&	6	F	V	f	v
7	BEL	ETB	'	7	G	W	g	w
8	BS	CAN	(	8	H	X	h	x
9	HT	EM	)	9	I	Y	i	y
A	LF	SUB	*	:	J	Z	j	z
B	VT	ESC	+	;	K	[	k	{
C	FF	FS	,	<	L	\	l	
D	CR	GS	-	=	M	]	m	}
E	SO	RS	.	>	N	^	n	~
F	SI	US	/	?	O	_	o	del





# UTF-8

- Unicode Transformation Format – 8-bit
- Variable length encoding
- Up to four bytes per symbol
- The first 128 are the same as for ASCII
  - Backwards compatibility – ASCII text is also valid UTF-8
- Dominating format on the Web



# There are Many Text/String Coding Schemes

- Families or series of encodings
  - EBCDIC (IBM)
- ISO 8859
- MS-Windows
- Mac OS Roman
- Unicode
  - > *UTF-8*
  - > *UTF-16*
  - ...

"knäckebröd av råg" encoded in UTF-8  
and decoded as ISO-8859-1 becomes  
"knÃæckebrÃ¶d av rÃ¥g"

*From Wikipedia*



# Encoding/Decoding in Java

- Methods and constructors to convert between strings and bytes
- Many take encoding scheme ("Charset") as parameter
  - > *Optional parameter*
  - > *If unspecified, there is a default encoding*
    - Typically UTF-8

```
String string = "Fruit flies like a banana";

// encode a string into a byte array
byte [] encodedBytes = string.getBytes(StandardCharsets.UTF_8);

// decode a byte array into a string
String decodedString = new String(encodedBytes, StandardCharsets.UTF_8);
```



# Other Data Types

- In general, any object needs a encoding/decoding scheme to export/import it in a program
  - Transfer over network, save in a file, ...
- Floating numbers
  - > *IEEE 754-2008*
  - > *Single-precision (approx 7 decimal digits): 32 bits*
  - > *Double precision (approx 16 decimal digits): 64 bits*
  - Integer
    - > *8, 16, 32, 64 bits (or more)*
    - > *Byte order – big-endian or little-endian?*
  - ...



# Complete Server with Data Processing

```
import java.net.*;
import java.nio.charset.StandardCharsets;

class TCPServer {
    static int BUFFERSIZE=1024;

    public static void main(String argv[]) throws Exception
    {
        ServerSocket welcomeSocket = new ServerSocket(6789);

        while(true) {
            Socket connectionSocket = welcomeSocket.accept();
            byte[] fromClientBuffer = new byte[BUFFERSIZE];
            int fromClientLength = connectionSocket.getInputStream().read(fromClientBuffer);

            String clientSentence = new String(fromClientBuffer, 0,
                                                fromClientLength, StandardCharsets.UTF_8);

            // Capitalize string
            String capitalizedSentence = clientSentence.toUpperCase();
            byte[] toClientBuffer = capitalizedSentence.getBytes(StandardCharsets.UTF_8);

            connectionSocket.getOutputStream().write(toClientBuffer);
            connectionSocket.close();
        }
    }
}
```

Class with Charset definitions for coding schemes

Decode byte sequence into string.

Encode string into byte sequence



# Java I/O Plumbing

- In the examples so far, we use explicit encoding and decoding
  - `String()` constructor, `getBytes()` method, etc
- In Java I/O, another method is to wrap readers/writers onto streams
- `InputStreamReader` converts a byte stream to a character stream by encoding it with a given character set
- `OutputStreamWriter` does the opposite

```
inStreamReader = new InputStreamReader(connectionSocket.getInputStream(),  
                                     StandardCharsets.UTF_8);  
  
outStreamWriter = new OutputStreamWriter(connectionSocket.getOutputStream(),  
                                          StandardCharsets.UTF_8);
```



# Java I/O Plumbing II

- With InputStream/OutputStream, every read/write operation leads to send/receive socket operations
- BufferedReader/BufferedWriter buffers data instead, and calls read/write when needed
  - For instance, when buffer gets full

```
inReader = new BufferedReader(new InputStreamReader(connectionSocket.getInputStream(),  
                                                    StandardCharsets.UTF_8));  
  
outWriter = new BufferedWriter(new OutputStreamWriter(connectionSocket.getOutputStream(),  
                                                       StandardCharsets.UTF_8));
```



# TCP Server According to Kurose-Ross, 5<sup>th</sup> ed.

```
import java.io.*;
import java.net.*;

class TCPServer {

    public static void main(String argv[]) throws Exception
    {
        String clientSentence;
        String capitalizedSentence;

        ServerSocket welcomeSocket = new ServerSocket(6789);

        while(true) {

            Socket connectionSocket = welcomeSocket.accept();

            BufferedReader inFromClient =
                new BufferedReader(new
                    InputStreamReader(connectionSocket.getInputStream()));
```

Create input  
stream attached  
to socket







# TCP Server According to Kurose-Ross 5<sup>th</sup> ed., cont

create output  
stream attached  
to socket



```
DataOutputStream outToClient =  
    new DataOutputStream(connectionSocket.getOutputStream());  
  
    clientSentence = inFromClient.readLine();  
  
    capitalizedSentence = clientSentence.toUpperCase() + '\n';  
  
    outToClient.writeBytes(capitalizedSentence);  
    }  
    }  
}
```



# Summary

- Socket Programming API
  - Create socket for communication
  - Perform operations on sockets (system calls)
    - > *connect()*, *bind()*, *listen()*, *read()*, *write()*, ...
- Byte transfer services
  - Sequence of bytes – stream (TCP)
  - Block of bytes – datagram (UDP)
- Conversion between data types in programming language and bytes
  - Encoding and decoding
- A variety of encoding schemes for conversion between characters (strings) and bytes
  - ASCII, UTF-8, ISO 8859-1, ...