IK1203

Dynamic Routing

Repetition: Basic Routing

- Our approach so far:
  - Next-hop routing
  - Logical (IP) addresses
  - Static Tables
- This approach works only for small IP networks
- We need to support dynamic large networks

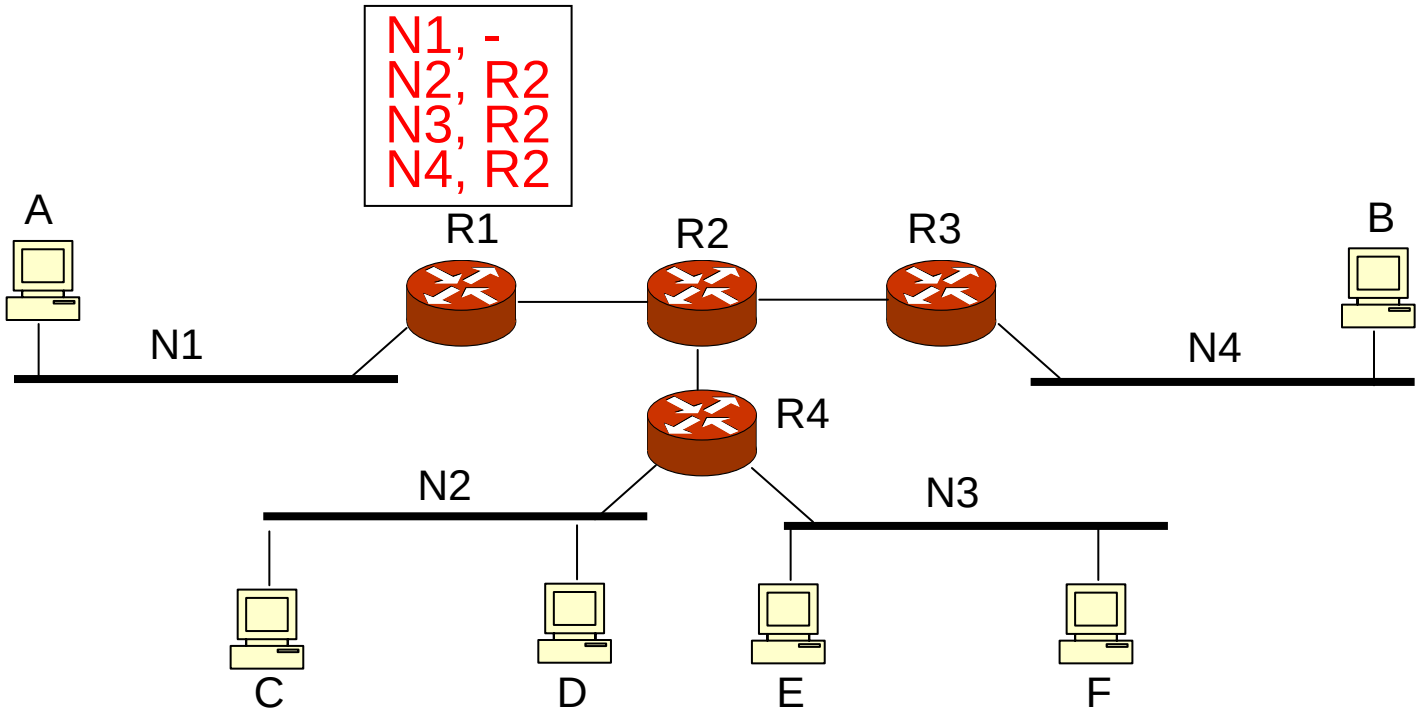

Levels of Abstraction

- The Internet is huge
  - Necessary to divide the routing problem into sub-problems.
  - There are several layers of *abstractions*
- Topmost: the Internet is partitioned into *Autonomous systems (AS)*
  - An independent administrative domain
  - Routing between AS:s is called inter-domain routing / External routing
  - Based on commercial agreements – Policies, Service-level-agreements
- Intermediate: An AS may be further partitioned into *areas*.
  - Routing inside an AS: Intra-domain routing / Internal routing
  - Best path based on hop/bw metrics
- Lowest level: Basic routing
  - Next-hop routing
  - Static routes

Partitioning: AS and Areas

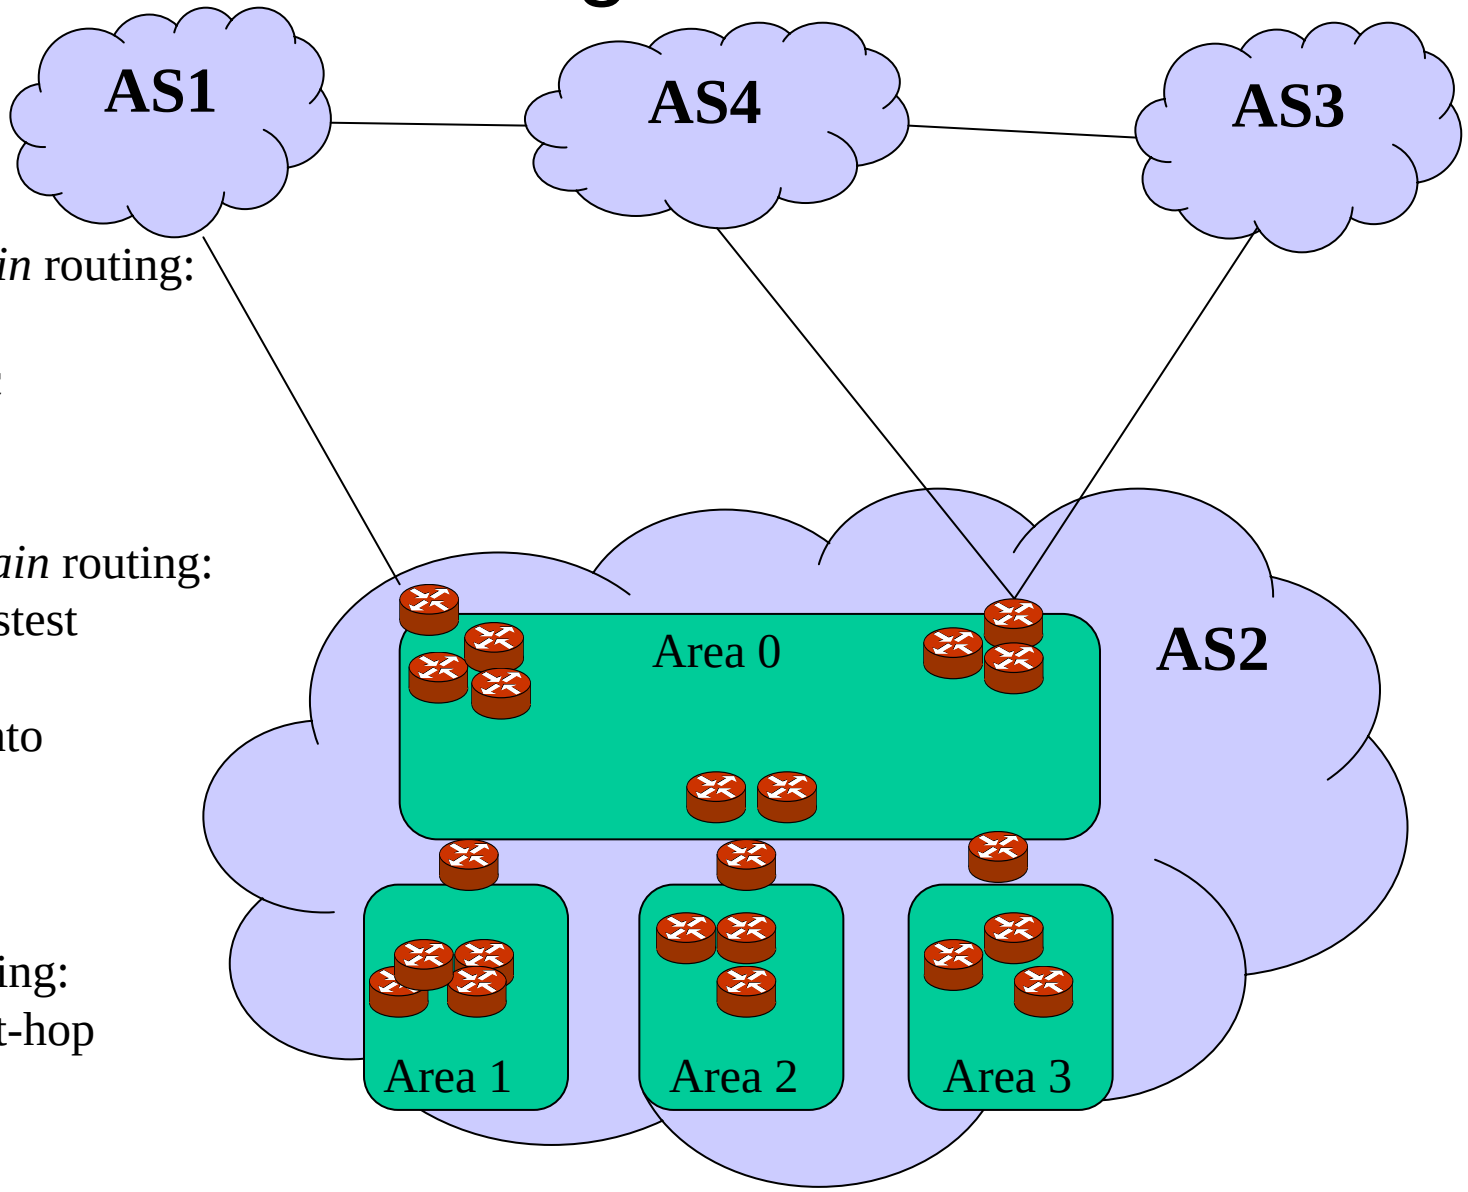

Inter-domain routing:
policies,
peering, etc

Intra-domain routing:
Finding fastest
path,
partition into
areas

Basic routing:
static, next-hop
routing

Autonomous Systems—RFC1930

- An *Autonomous system* is generally administered by a single entity.
  - Operators, ISPs (Internet Service Providers)
- An AS contains an arbitrary complex sub-structure.
- Each autonomous system selects the routing protocol to be used *within* the AS.
- Policies or updates within an AS are not propagated to other AS:s.
- An AS-number is (currently) a 16-bit unique identifier
- Interconnwection between AS:s
  - Service Level Agreements (SLA:s)
  - Internet Exchange Points (IX:s)
  - Network Access Points (NAPs)

AS Number	Network
3	MIT
32	STANFORD
2839	KTH
1653	SUNET

Sample Internet Architecture

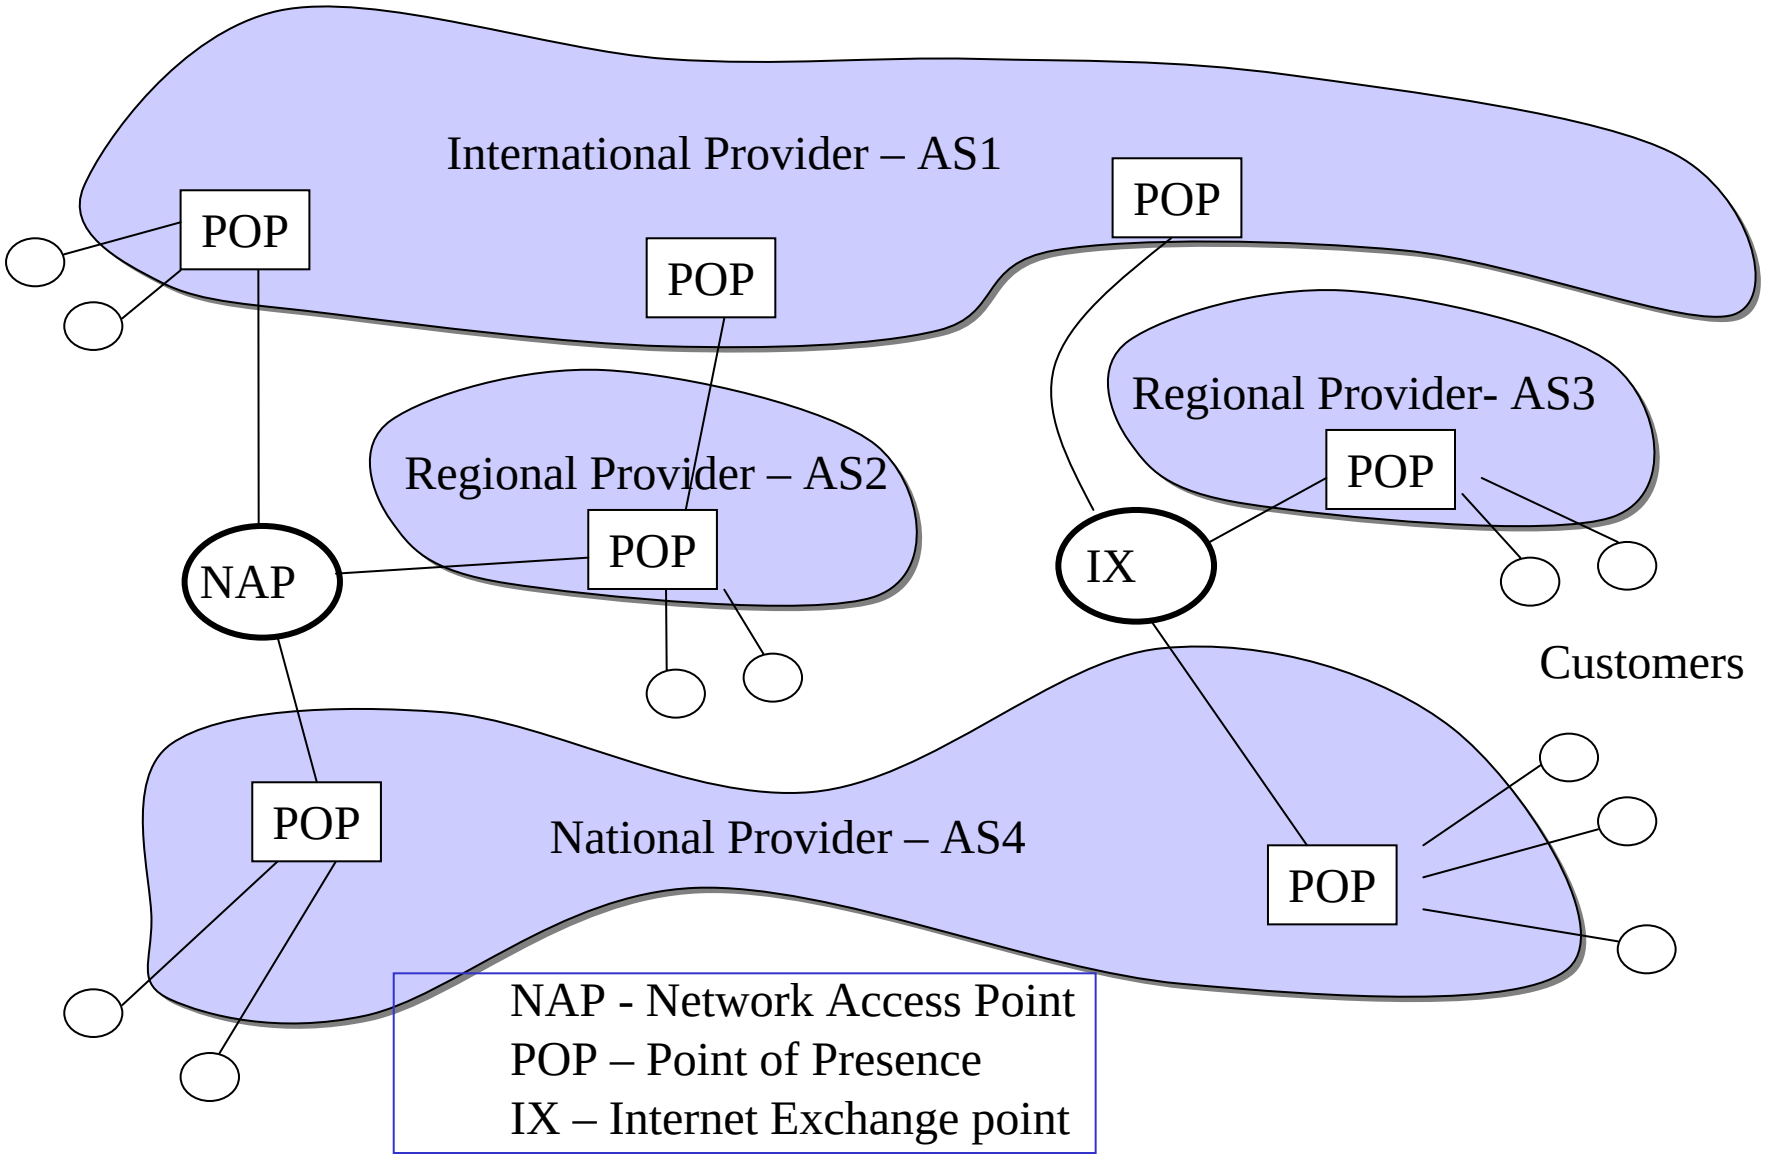

Reachability and Metrics

- The most fundamental functionality in a dynamic routing protocol:
  - Find the ”best path” to a destination
- Two algorithms in use to find best path
  - Distance-Vector (Bellman-Ford)
  - Link-state (Dijkstra)
- But what is best path?
  - Interior routing: typically number of hops, or bandwidth
  - Exterior routing: business relations—peering
- Metrics
  - Number of hops (most common)
  - Bandwidth, Delay, Cost, Load, ”Policies”

Aggregation

- Also called *summarization*
- The netid part of IPv4 addresses can be aggregated (summarized) into shorter prefixes.
  - Currently: ~160000 global prefixes
- Summarization is often done manually
- Leads to smaller routing tables (fewer prefixes)
- Threats: multi-homing and load-balancing

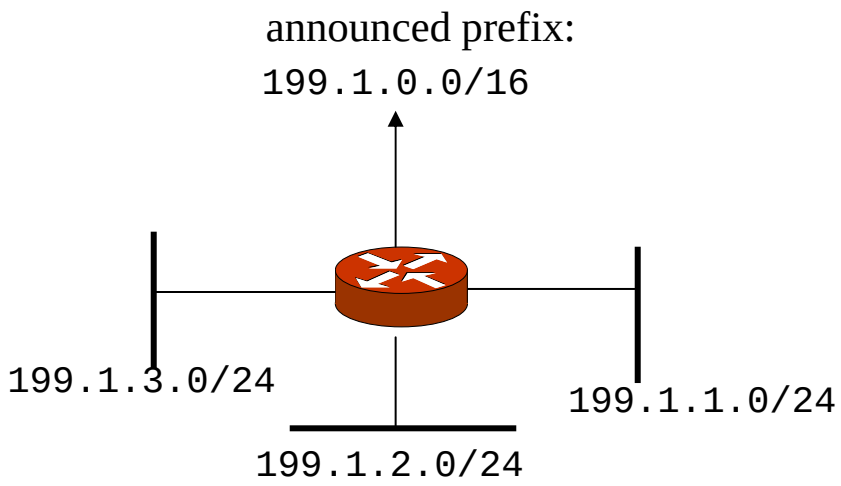

Redistribution of Routing Information

- If several protocols are running on the same router
  - E.g., an OSPF as interior and BGP as exterior
  - E.g. static routes into dynamic routing protocol
- The router can distribute routes from one protocol to another
  - Interior routes need to be advertized to the Internet
    - Typically these routes are aggregated
  - Exterior routes may need to be injected into the interior network
    - But only a subset—the backbone tables are very large
    - Necessary for domain carrying *transit* traffic
    - Not necessary for a domain using only a default route
- Typically, redistributed routes are filtered in different ways

Load Balancing

- The routing protocol gives several routes to a network
- Either select the best
- Or load-balance between several links
  - Unequal-cost multi-path
  - Equal-cost multi-path
- The forwarding decides *how* to balance actual traffic:
  - Randomly (not good for TCP)
  - Load balance per flow
  - Load balance per address pairs

Popular Routing Protocols

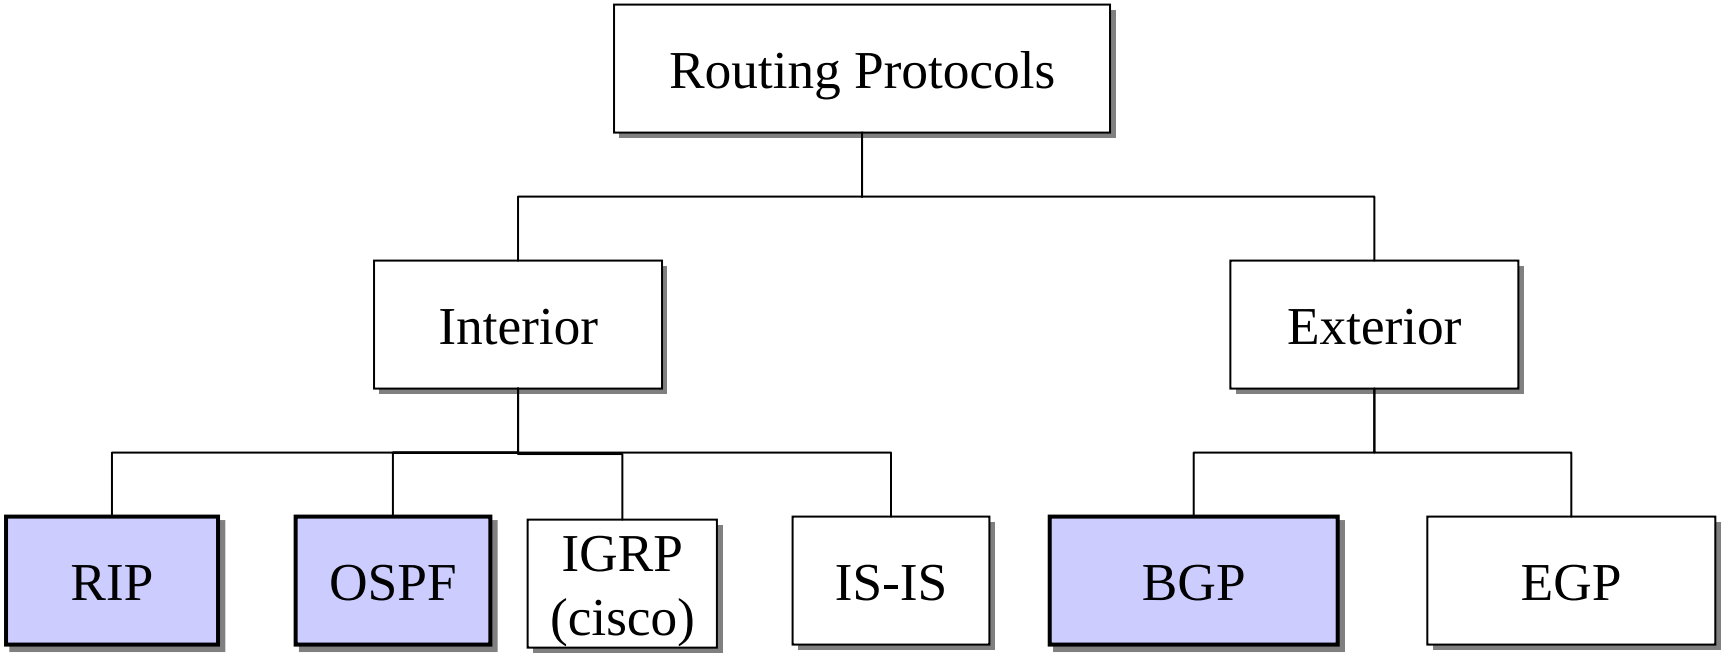

Routing Information Protocol - RIP

- RIP-1 (RFC 1058), RIP-2 (RFC 2453)
- Metric is Hop Counts
  - 1: directly connected
  - 16: infinity
  - RIP cannot support networks with diameter > 15.
- RIP uses distance vector
  - RIP messages contain a vector of hop counts.
  - Every node sends its routes to its neighbours
  - Route information gradially spreads through the network
  - Every node selects the route with smallest metric.
- RIP messages are carried via UDP datagrams.
  - IP Multicast (RIP-2) or Broadcast (RIP-1)

Distance Vector

- A node advertizes its “distance-vector”
  - A list (vector) of all nodes that the node knows about
  - The distance to each of them
- Advertizements are sent to neighbours only
- Each neighbour updates its routing table and sends the new distance-vectors to its neighbours
    - Bellman-Ford algorithm

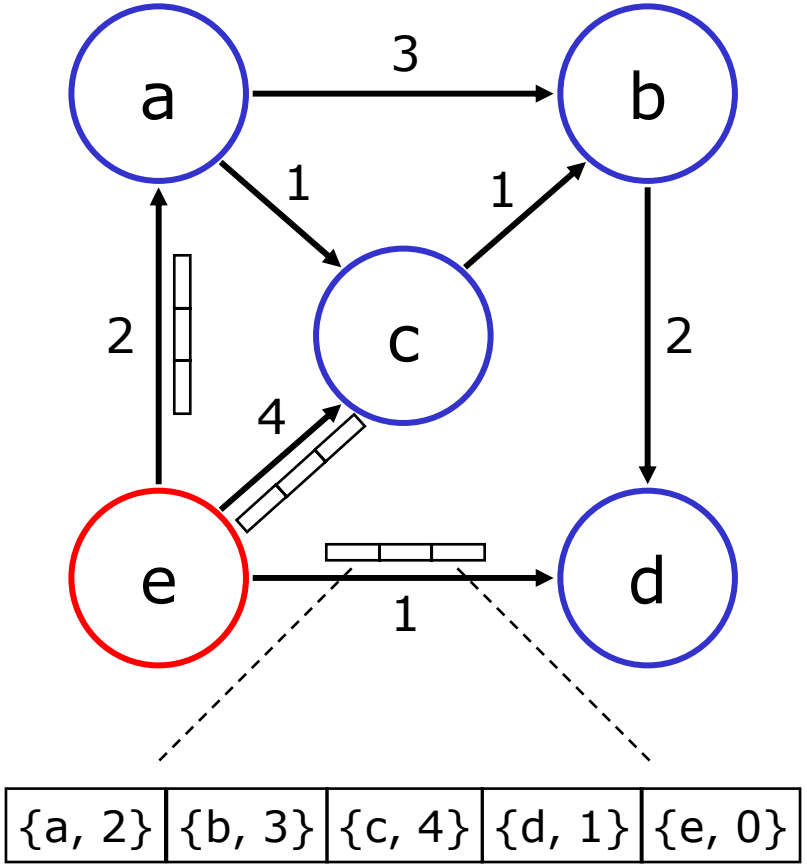

Distance-vector from ”e”

RIP Problem: Count to Infinity

1. Initially, R_1 and R_2 both have a route to N with metric 1 and 2, respectively.

2. The link between R_1 and N fails.

3. Now R_1 removes its route to N, by setting its metric to 16 (infinity).

4. Now two things can happen: Either R_1 reports its route to R_2. Everything is fine.

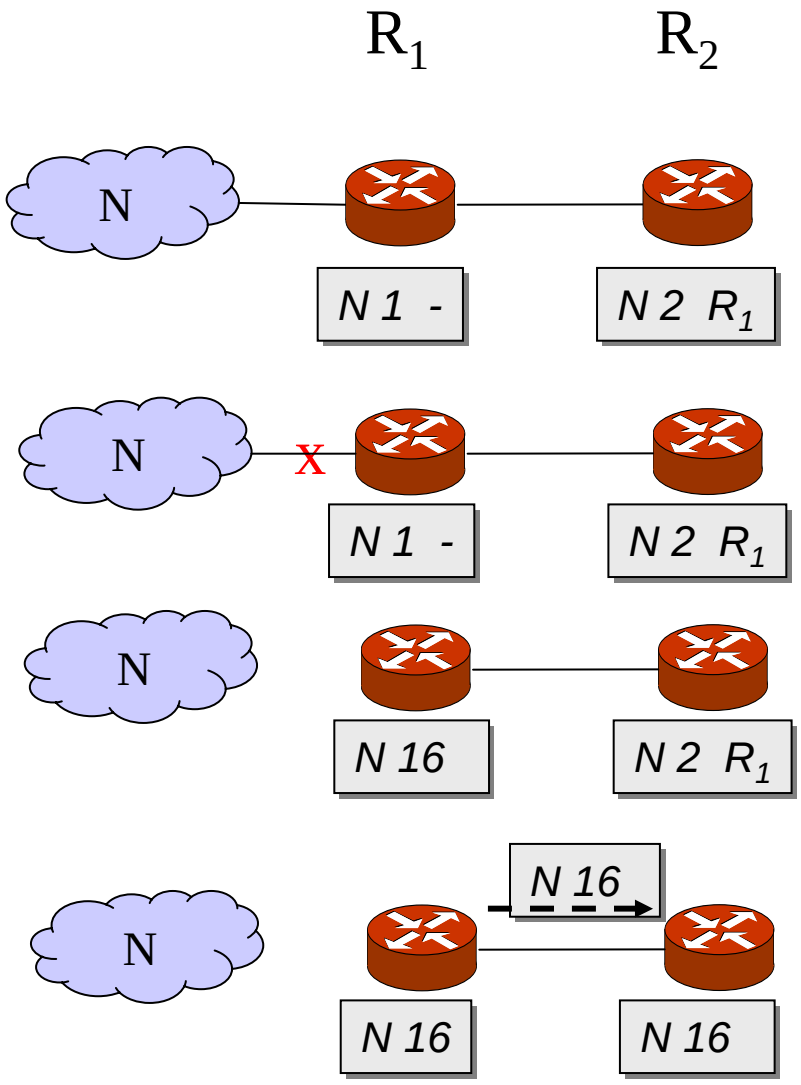

RIP Problem: Count to Infinity

5. The other alternative is that R_2, which still has a route to N, advertises it to R_1. Now things start to go wrong: packets to N are looped until their TTL expires!

6. Eventually (~10-20s), R_1 sends an update to R_2. The cost to N increases, but the loop remains.

7. Yet some time later, R_2 sends an update to R_1.

   ...

13. Finally, the cost reaches infinity at 16, and N is unreachable. The loop is broken!

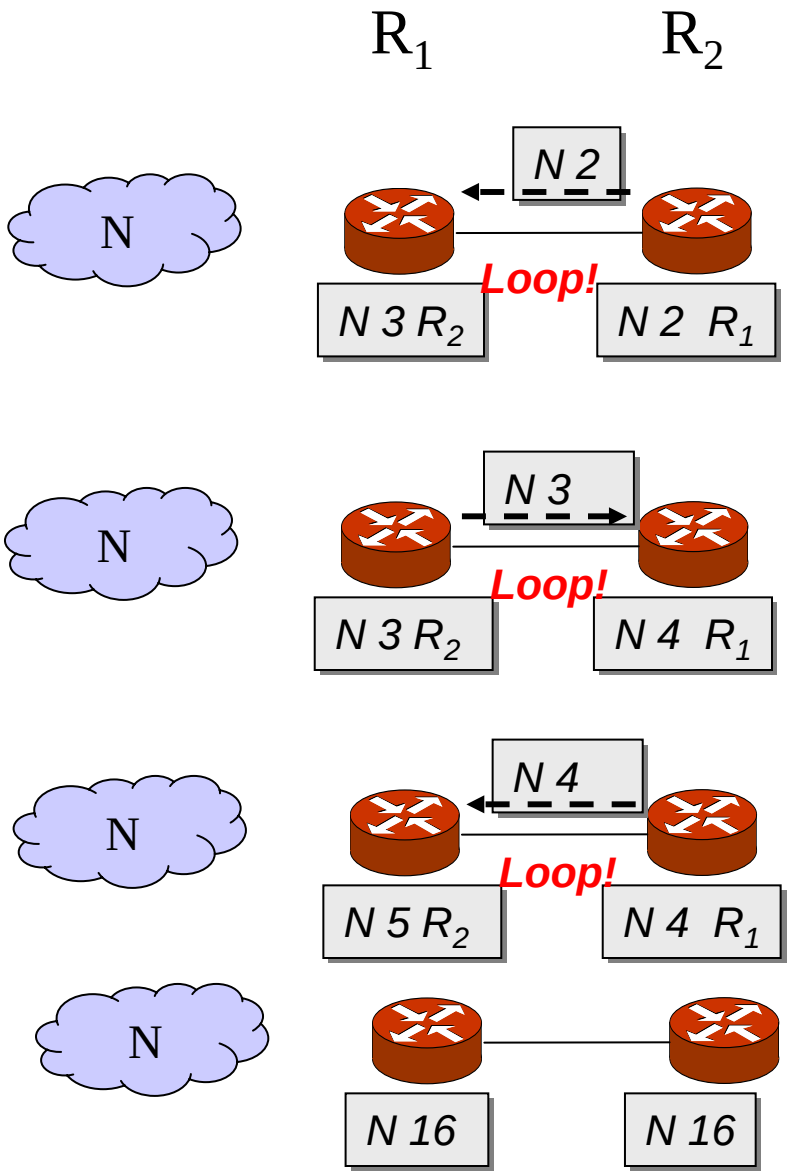

Solution1: Triggered Update

- Send out update immediately

R_1 Immediately announces the broken link when it happens.

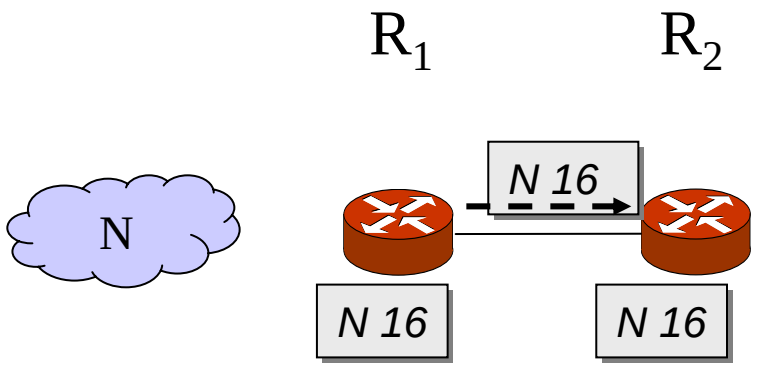

Solution 2: Split Horizon

- Do not propagate information about a route over the same interface from which the route arrived.
- Split horizon only prevents loops between adjacent routers

R_2, does not announce the route to N to R_1 since that is where it came from.

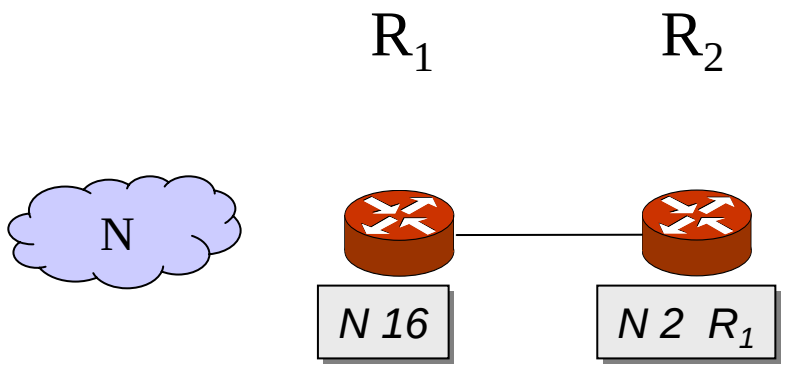

Eventually, R_1 reports its route to R_2 and everything is fine.

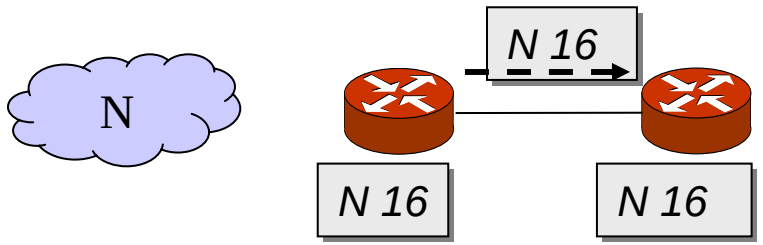

Split Horizon Does Not Solve All Cases

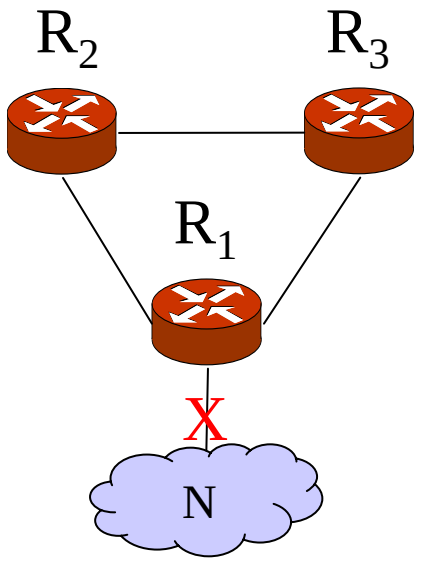

- R1 reports to R2 that N is unreachable
- R2 believes that D is reachable through R3
- R2 reports this to R1, who believes that N can be reached through R2....

Somewhat depending on timing.....

Solution 3: Poison Reverse

- Advertise reverse routes with a metric of 16 (i.e., unreachable).
- Somewhat more aggressive variant of split horizon
  - It handles some more error cases than split horizon

R_2 always announces an unreachable route to N to R_1.

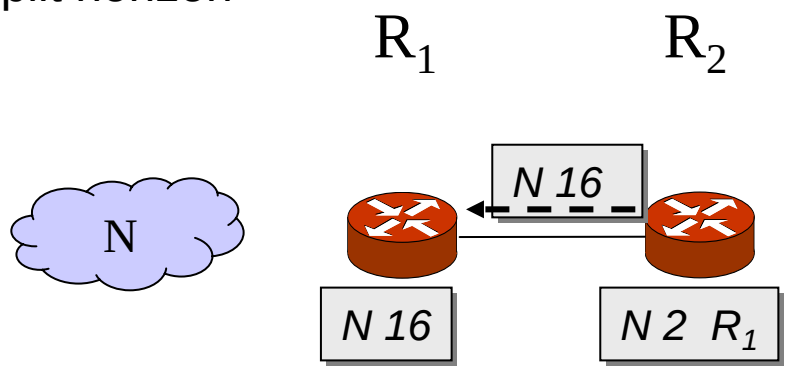

Eventually, R_1 reports its route to R_2 and everything is fine.

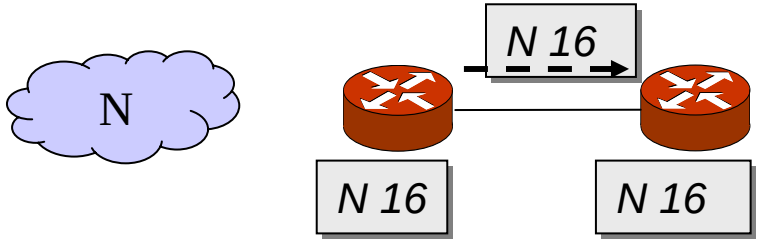

Solution 4: Hold Down

- When a route is removed, no update of this route is accepted for some period of time (hold-down time)- to give everyone a chance to remove the route.

R_1 ignores updates to N from R_2 for some period of time.

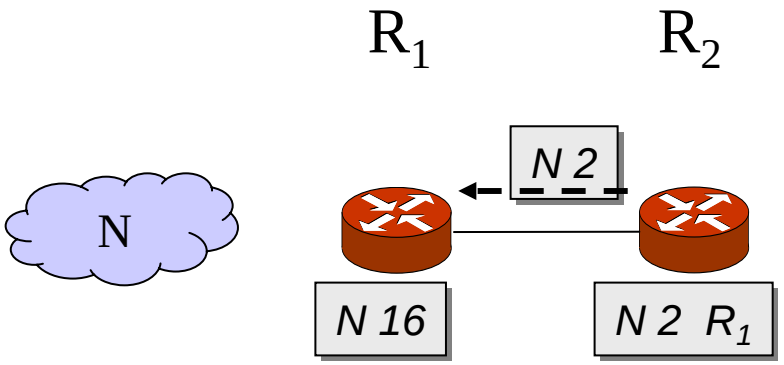

Eventually, R_1 sends the update to R_2.

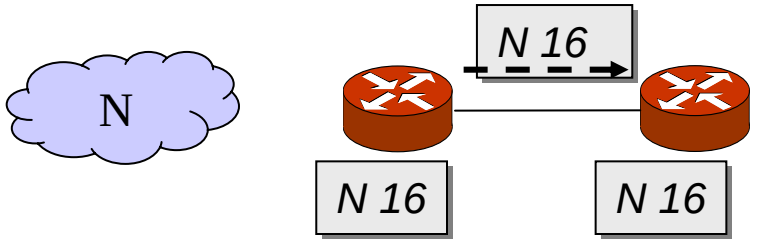

Disadvantages with RIP

- Slow convergence
  - Changes propagate slowly
  - Each neighbor only speaks ~every 30 seconds; information propagation time over several hops is long
- Instability
  - After a router or link failure RIP takes *minutes* to stabilize.
- Hops count may not be the best indication for which is the best route.
- The maximum useful metric value is 15
  - Network diameter must be less than or equal to 15.
- RIP uses lots of bandwidth
  - It sends the whole routing table in updates.

Why Use RIP?

- After all these problems you might ask this question
- Answer
  - Because RIP is generally available
  - It is simple to configure.

Open Shortest Path First—OSPF

- OSPF version 2
  - RFC 2328
- OSPF is a link-state protocol.
  - Builds *Link State Advertisements* (LSAs)
  - Distributes LSAs to all other routers
  - Computes delivery tree using the *Dijkstra* algorithm
- OSPF uses IP *directly* (protocol field = 89)
  - Not UDP or TCP.
- OSPF networks are partitioned into *areas* to minimize cross-area communication.

OSPF Network Topology

- Area 0 is the *backbone* area. All traffic goes via the backbone.
- All other areas are connected to the backbone (1-level hierarchy)
- *A Border area router* has one interface in each area.
- An *AS Boundary Router*—attaches to other AS:s

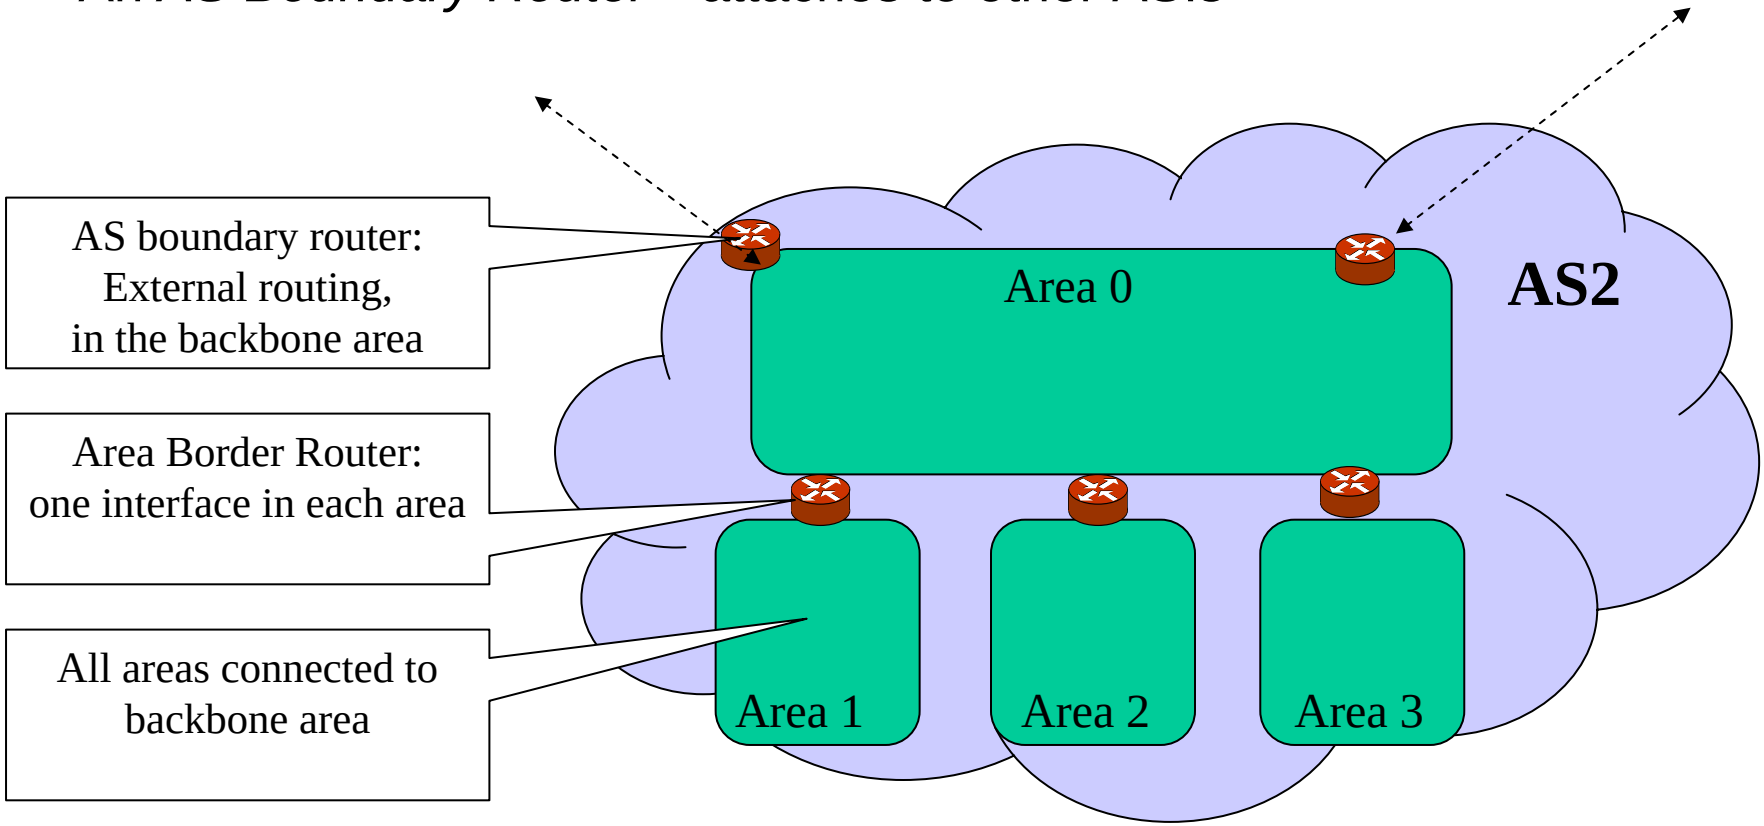

Link-State Protocols (SPF)

- In SPF, every router does the following:
  1. Actively test the status of all neighbours/links
  2. Build a Link State Advertisement (LSA) from this information and propagate it to *all other* routers within an area.
  3. Using LSAs from all other routers, compute a shortest path delivery tree, typically using *Dijkstra shortest path algorithm*.
- Advantages (over distance-vector):
  - More functionality due to computation on original data and no dependence on intermediate routers
    - Full topology knowledge
    - Easier to Troubleshooting
  - Fast Convergence
- Disadvantage
  - uses more memory

OSPF Contains Three Protocols

1. The *Hello* protocol
   - Check for neighbours, authentication, designated routers
2. The *Exchange* Protocol
   - Exchange Link State Database between neighbours
   - First get LSA headers
   - Then transfer actual LSAs on request.
3. The *Flooding* protocol
   - When links change/age
   - Send Link State updates to neighbours and flood *recursively*.
   - If not seen before, propagate updates to all *adjacent* routers, except incoming

Distribution of Link State Advertisments

- Most complex and critical part of OSPF
- Initial topology transfer done with the exchange protocol.
- OSPF *floods* LSAs within an *area*
  - Recursively forward a new LSA to all neighbours (except the recepient)
  - An LSA will travel on all links exactly once
  - Uses sequence numbers and aging to avoid loops
- OSPF aggregates routes
  - Border Area Routers aggregates routes from an area into other areas.
  - AS Border Routers aggregates routes from other ASs.

Dijkstra Algorithm (Shortest Path First)

Find shortest paths from "a" to all other nodes!

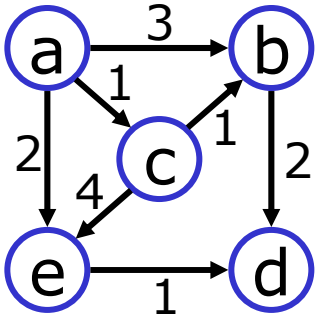

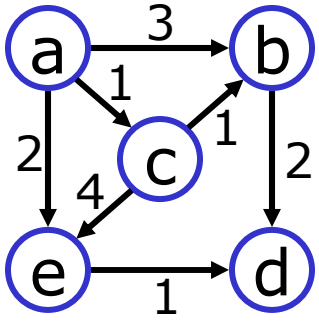

M	D_b (path)	D_c (path)	D_d (path)	D_e (path)
{a}	3 (a-b)	1 (a-c)	∞ (--)	2 (a-e)
{a, c}	2 (a-c-b)	**1 (a-c)**	∞ (--)	2 (a-e)
{a, c, b}	**2 (a-c-b)**	**1 (a-c)**	4 (a-c-b-d)	2 (a-e)
{a, c, b, e}	**2 (a-c-b)**	**1 (a-c)**	3 (a-e-d)	**2 (a-e)**
{a, c, b, e, d}	**2 (a-c-b)**	**1 (a-c)**	**3 (a-e-d)**	**2 (a-e)**

Alternative to OSPF: IS-IS

- Link-State Routing
- Originally designed for Decnet and then CLNP (OSI)
- Has been stable for a longer time than OSPF
  - Large deployed base
  - Example: SUNET runs IS-IS
- More general hierarchies
  - Multiple levels in tree topology
  - Not strict two-levels as OSPF

Border Gateway Protocol—BGP

- Inter-domain routing
- Simple cases: *use static routing*
- Main purpose: Network reachability between autonomous systems
- BGP version 4 is *the* exterior routing protocol used in the Internet today.
- BGP uses TCP
  - TCP is reliable: reduces the protocol complexity
- BGP uses *path-vector* - enhancent of distance-vector.
- BGP implements *policies* – chosen by the local administrator.

BGP Simple example

- AS1 has a network 1.0.0.0/8 that it announces

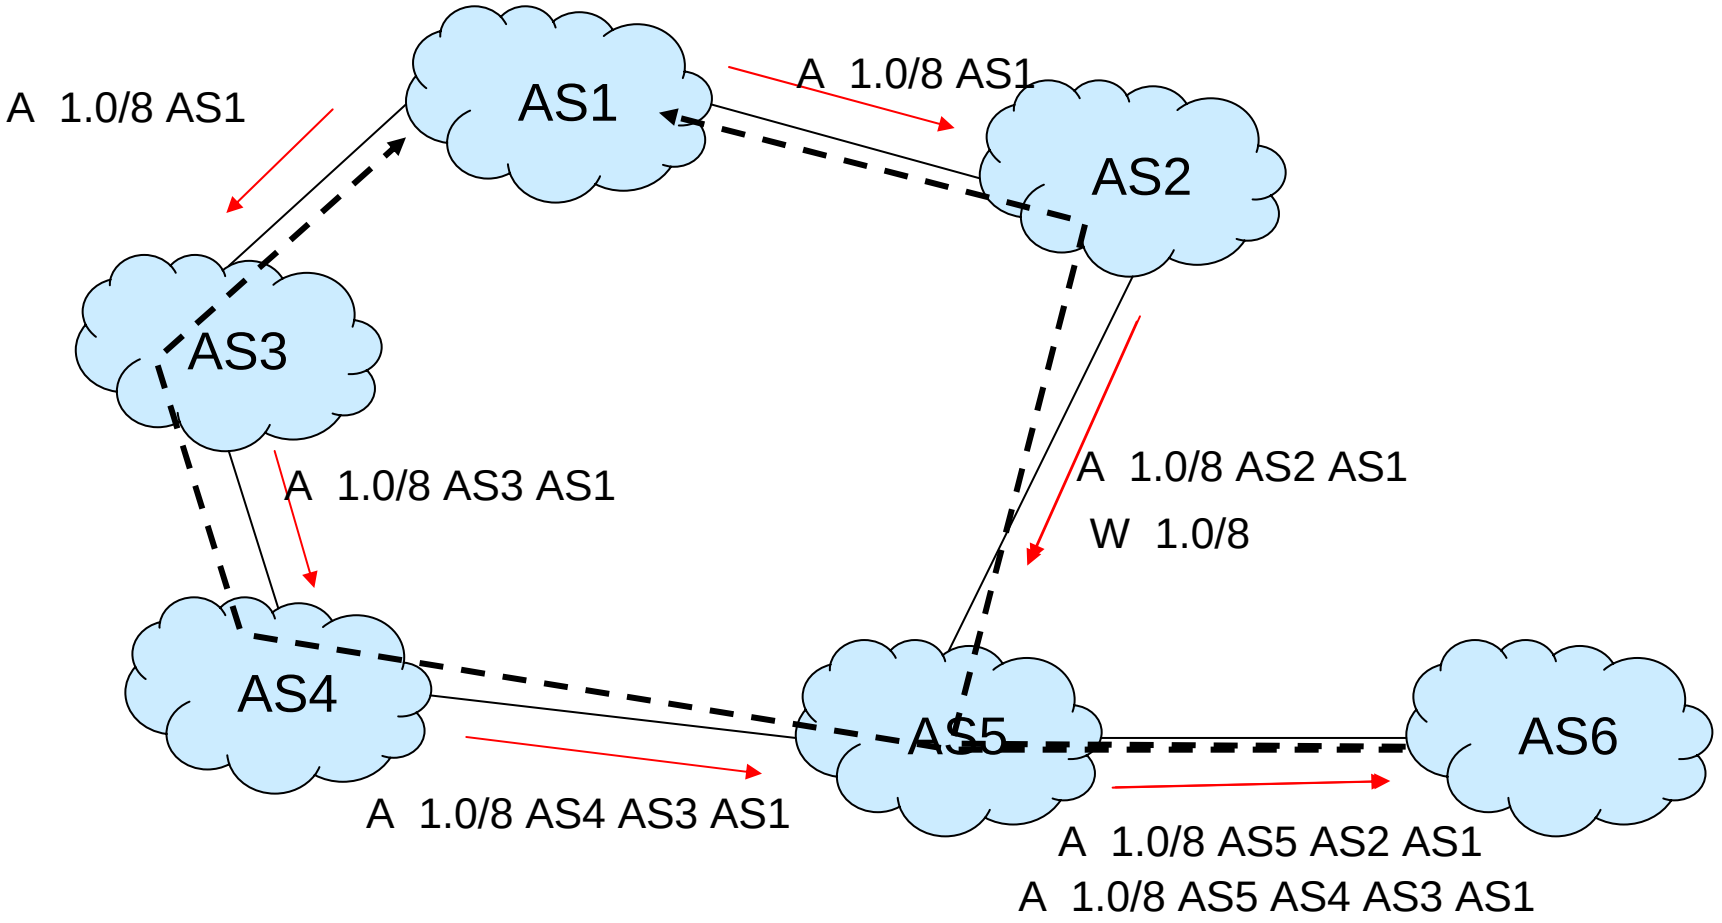

Motivation for Path-Vector

- Distance-vector
  - Hop-count too limited
  - Unstable
- Link-State
  - Link state database would be enormous
- Path-vector extends distance-vector
  - Instead of a simple cost, assign *an AS-Path* to every route
  - There may be many paths to the same destination (network *prefix*)
  - AS-Path used to implement *policies* and *loop prevention*

BGP Architecture

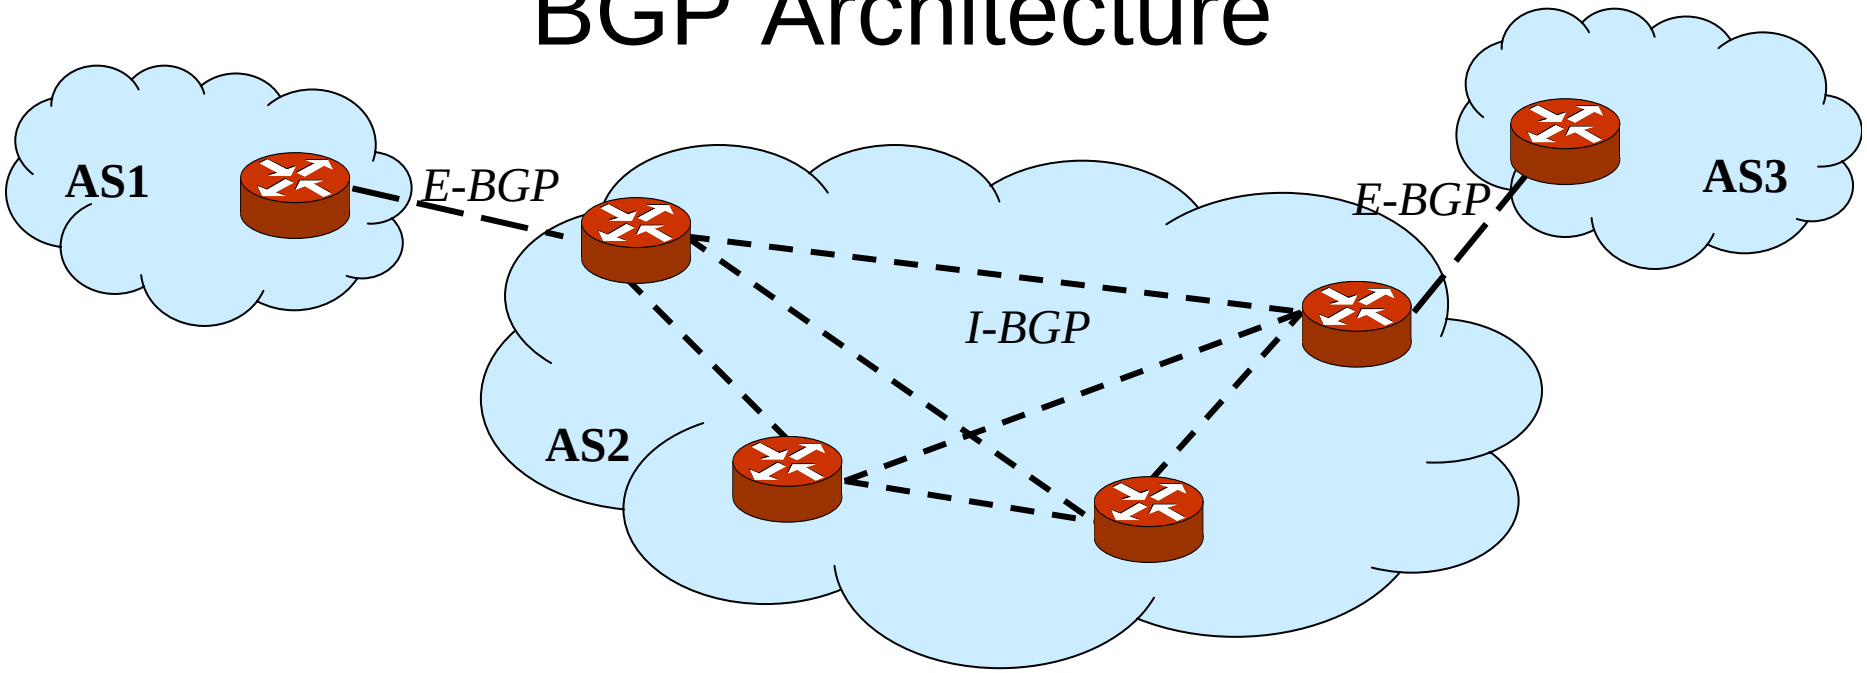

- BGP interacts with the internal routing (OSPF/IS-IS/RIP/...)
  - Redistributes routes between the two domains
- BGP really consists of two protocols:
  - E-BGP: coordinates between border routers *between* AS:s
  - I-BGP : coordinates between BGP peers *within* an AS

BGP Router Operation

- A BGP router receives routes
  - BGP peers (E-BGP)
  - Redistribution: IGP/static routes
- It aggregates routes
- It filters and modifies routes
  - According to some *policy*
- It advertizes routes to its EBGP neighbours in other AS:s

BGP Router Model

- One adjacent-in-RIB per peer
- Selects routes to use (policy-based)
- Advertizes routes to peers: One adjacent-out-RIB per peer

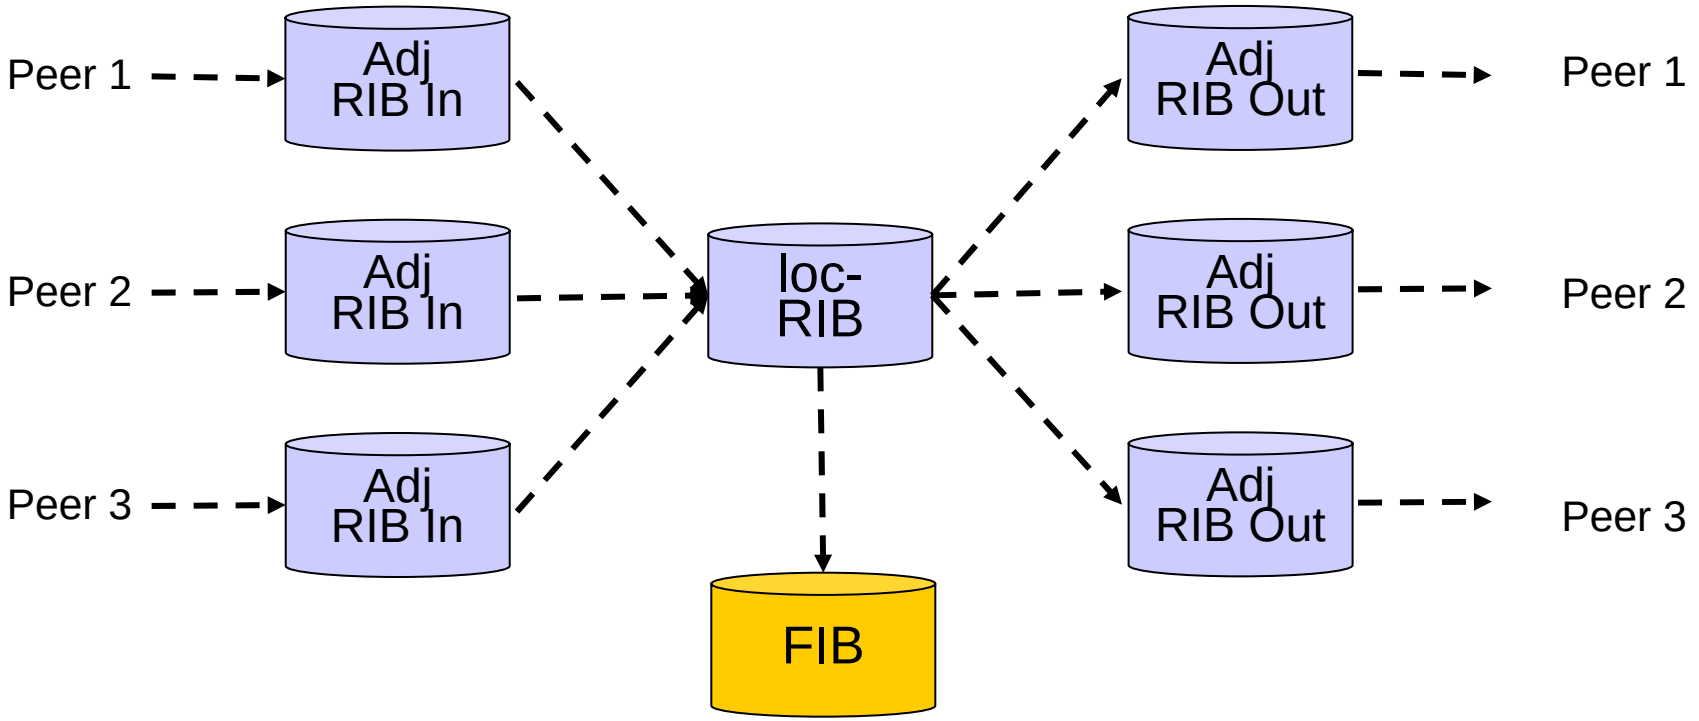

External BGP scaling

- Currently , the problems BGP faces have to do with large and changing routing tables
- Routing tables are large because of the many subnetworks (~160000)
- Aggregation does not always work
  - E.g., Multihoming often breaks aggregation
- When networks change, BGP peers advertise changes, and these may have problems to converge
  - Especially in the presence of faults
- Valid BGP implementations are few, and it is difficult to configure BGP

Internal BGP scaling

- External routes need to be distributed within an AS: I-BGP
- But all I-BGP peers need to be connected in a full mesh
  - This does not scale
- Route reflection
  - Break up the mesh into a hierarchy
- AS confederations
  - divide-and-conquer
  - use several sub-ASs